

transwarp

A header-only C++ library for task concurrency


Christian Blume

May 17, 2017

transwarp drive

Subspace short cuts

Transwarp corridors are effectively subspace short cuts between far distant sections of the Galaxy. By using these corridors, the Borg are able to travel hundreds of light years in a matter of minutes, meaning that they can traverse the massive distance between the Delta and Alpha Quadrants in a relatively short period of time.



transwarp library

A header-only C++ library for task concurrency

<https://github.com/bloomen/transwarp>

- ▶ Written in C++11
- ▶ Currently in version 0.1.0
- ▶ Tested using GCC and Clang
- ▶ MIT license

Let's chat if you want to contribute!

The problem

- ▶ we have a bunch of operations that depend on each other
- ▶ those operations are run repeatedly with varying input
- ▶ some of those operations are independent

Now, we want to:

- ▶ understand the dependencies between operations
- ▶ run the independent operations in parallel

Typical solution

- ▶ launch a thread per operation that should run concurrently
- ▶ somehow notify a given operation that it is ready to run
- ▶ collect the results and move on to the next operation


Drawbacks

- ▶ explicit thread handling
- ▶ a lot of independent operations may lead to thread contention
- ▶ error prone notification mechanism
- ▶ dependencies are hard to reason about

We can do better!

Solving it with transwarp

- ▶ directed acyclic graph of tasks
- ▶ a task can have any number of parent tasks
- ▶ graph must contain a single final task without children
- ▶ underlying thread pool to actually run the tasks
- ▶ a task operation may run anywhere (CPU, GPU, etc)

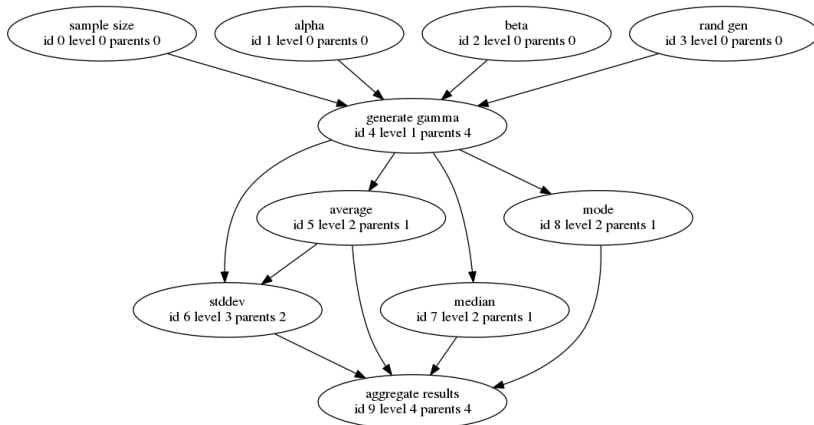


Simple example

```
1  double add_em_up(double x, int y) {
2      return x + y;
3  }
4
5  int main() {
6      double value1 = 13.3;
7      int value2 = 42;
8      auto something = [&value1] { return value1; };
9      auto something_else = [&value2] { return value2; };
10
11     auto t1 = make_task("something", something);
12     auto t2 = make_task("something else", something_else);
13     auto t3 = make_final_task("adder", parallel{4},
14                               add_em_up, t1, t2);
15
16     t3->schedule();
17     std::cout << t3->get_future().get(); // 55.3
18 }
```

A slightly more evolved example

- ▶ generate data from a gamma distribution
- ▶ compute statistical measures such as average and median
- ▶ check out the code on github



The task class

```
1  template<typename Functor, typename... Tasks>
2  class task {
3  public:
4
5      task(std::string name, Functor functor,
6           std::shared_ptr<Tasks>... parents);
7
8      std::shared_future<result_type> get_future() const;
9
10     const node& get_node() const;
11
12     template<typename Pre, typename Post>
13     void visit(Pre& pre_visitor, Post& post_visitor);
14
15     void unvisit();
16 };
```

Visiting tasks

```
1  template<typename Pre, typename Post>
2  void task::visit(Pre& pre, Post& post) {
3      if (!visited_) {
4          pre(*this);
5          detail::visit<Pre, Post> visit_task(pre, post);
6          detail::call_with_each(visit_task, parents_);
7          post(*this);
8          visited_ = true;
9      }
10 }
11
12 void task::unvisit() {
13     if (visited_) {
14         visited_ = false;
15         detail::call_with_each(detail::unvisit(), parents_);
16     }
17 }
```

Supporting classes

```
1  struct node {
2      std::size_t id;
3      std::size_t level;
4      std::string name;
5      std::vector<const node*> parents;
6  };
7
8
9
10 struct edge {
11     const transwarp::node* child;
12     const transwarp::node* parent;
13 };
```

The final task class

```
1  template<typename Functor, typename... Tasks>
2  class final_task : public task<Functor, Tasks...> {
3  public:
4
5      final_task(std::string name, sequenced seq, Functor f,
6                  std::shared_ptr<Tasks>... parents);
7
8      final_task(std::string name, parallel par, Functor f,
9                  std::shared_ptr<Tasks>... parents);
10
11     void schedule();
12
13     void set_pause(bool enabled);
14
15     void set_cancel(bool enabled);
16
17     const std::vector<edge>& get_graph() const;
18 };
19
20 std::string make_dot(const std::vector<edge>& graph);
```

Execution policies

```
1  class sequenced {};  
2  
3  
4  
5  class parallel {  
6  public:  
7      parallel(std::size_t n_threads);  
8  
9      std::size_t n_threads() const;  
10 };
```

Scheduling tasks

```
1 void final_task::schedule() {
2     if (!*canceled_) {
3         prepare_callbacks();
4         if (pool_) { // parallel execution
5             for (const auto& callback : callbacks_) {
6                 pool_->push(callback);
7             }
8         } else { // sequential execution
9             for (const auto& callback : callbacks_) {
10                 while (paused_) {};
11                 callback();
12             }
13         }
14     }
15 }
```

Creating packagers

```
1  detail::wrapped_packager task::make_packager() {
2      auto packager = [this] {
3          auto futures = detail::get_futures(parents_);
4          auto pack_task = std::make_shared<
5              std::packaged_task<result_type()>>(std::bind(
6                  &task::evaluate, std::ref(*this), futures));
7          future_ = pack_task->get_future();
8          return [pack_task] { (*pack_task)(); };
9      };
10     return {packager, &node_};
11 }
12
13 static
14 result_type evaluate(task& t, std::tuple<...> futures) {
15     if (*t.canceled_)
16         throw task_canceled(t.get_node());
17     return detail::call_with_futures<result_type>(
18         t.functor_, futures);
19 }
```

Scheduling tasks (recap)

```
1 void final_task::schedule() {
2     if (!*canceled_) {
3         prepare_callbacks();
4         if (pool_) { // parallel execution
5             for (const auto& callback : callbacks_) {
6                 pool_->push(callback);
7             }
8         } else { // sequential execution
9             for (const auto& callback : callbacks_) {
10                 while (paused_) {};
11                 callback();
12             }
13         }
14     }
15 }
```


The thread pool

```
1  std::queue<std::function<void()>> functors_;
2  std::condition_variable cond_var_;
3  std::mutex mutex_;
4
5
6  void thread_pool::push(std::function<void()> func) {
7      {
8          std::lock_guard<std::mutex> lock(mutex_);
9          functors_.push(func);
10     }
11     cond_var_.notify_one();
12 }
```

The thread pool (cont.)

```
1 void thread_pool::worker() {
2     for (;;) {
3         std::function<void()> func;
4         {
5             std::unique_lock<std::mutex> lock(mutex_);
6             cond_var_.wait(lock, [this]{
7                 return !paused_ && (done_ || !functors_.empty());
8             });
9             if (done_ && functors_.empty())
10                 break;
11             func = functors_.front();
12             functors_.pop();
13         }
14         func();
15     }
16 }
```

Summary

Use transwarp when

- ▶ you have dependent operations
- ▶ the operations are executed repeatedly
- ▶ you consider to run certain operations in parallel

What you get is

- ▶ parallelism by design
- ▶ a task graph that is easy to reason about
- ▶ focus on getting stuff done

Outlook

- ▶ consider notifications for when futures have been replaced
- ▶ test with C++17 enabled
- ▶ test with Visual Studio
- ▶ add some more real-life examples
- ▶ lift the library to version 1.0
- ▶ ...

Thank You